

## Table of contents

- [Table of contents](#)
- [S3 events](#)
- [Cloudwatch custom metrics :](#)
  - [Creating memory utilization metric](#)
  - [Auto scaling](#)
    - [Configuring auto scaling group](#)

## S3 events

- S3 events is similar to cloudwatch events . This is commonly used for object level activities like put,get etc
- S3 events sends out notification to SNS,lamba and SQS when the action is called.
- In order to use the SNS for S3 events however , we need to add access policy to the SNS topic.
- Navigate to SNS and click on edit the acces policy

```
{
  "Version": "2012-10-17",
  "Id": "__default_policy_ID",
  "Statement": [
    {
      "Sid": "__default_statement_ID",
      "Effect": "Allow",
      "Principal": {
        "AWS": "*"
      },
      "Action": [
        "SNS:GetTopicAttributes",
        "SNS:SetTopicAttributes",
        "SNS:AddPermission",
        "SNS:RemovePermission",
        "SNS:DeleteTopic",
        "SNS:Subscribe",
        "SNS:ListSubscriptionsByTopic",
        "SNS:Publish",
        "SNS:Receive"
      ],
      "Resource": "arn:aws:sns:us-east-1:384395217903:CW-alarm",
      "Condition": {
        "StringEquals": {
          "AWS:SourceOwner": "384395217903"
        }
      }
    },
    {
      "Sid": "bucketaccess",
      "Effect": "Allow",
      "Principal": {
        "Service": "s3.amazonaws.com"
      },
      "Action": [
        "s3:PutObject",
        "s3:GetObject",
        "s3:DeleteObject"
      ],
      "Resource": "*"
    }
  ]
}
```

```

    },
    "Action": "SNS:Publish",
    "Resource": "arn:aws:sns:us-east-1:384395217903:CW-alarm",
    "Condition": {
      "ArnLike": {
        "AWS:SourceArn": "arn:aws:s3:::aws-devops-4"
      }
    }
  }
]
}

```

- Use policy as above, just make sure to replace the account id in source owner along with topic arn and bucket arn
- Once done, navigate to S3. Go to properties and click on s3 events
- Click on add notifications. Give it an appropriate name and select the event on which we want notification for ex put
- Select the sns topic and click on save

**Name** ⓘ

test

**Events** ⓘ

<input checked="" type="checkbox"/> PUT	<input type="checkbox"/> All object delete events
<input type="checkbox"/> POST	<input type="checkbox"/> Restore initiated
<input type="checkbox"/> COPY	<input type="checkbox"/> Restore completed
<input type="checkbox"/> Multipart upload completed	<input type="checkbox"/> Replication time missed threshold
<input type="checkbox"/> All object create events	<input type="checkbox"/> Replication time completed after threshold
<input type="checkbox"/> Object in RRS lost	<input type="checkbox"/> Replication time not tracked
<input type="checkbox"/> Permanently deleted	<input type="checkbox"/> Replication failed
<input type="checkbox"/> Delete marker created	

**Prefix** ⓘ

e.g. images/

**Suffix** ⓘ

e.g. .jpg

**Send to** ⓘ

SNS Topic

**SNS**

CW-alarm

Disabled

- Test the event by uploading an object in S3

## Cloudwatch custom metrics :

- We have seen that cloudwatch gives us certain pre-defined metrics.
- Along with the above, we can also define some custom metrics.
- For example, even though we have CPU utilization metric by default, we don't have memory (RAM) utilization metric

### Creating memory utilization metric

- Launch an ec2 instance

- With some versions of Linux, you must install additional Perl modules before you can use the monitoring scripts. For amazon linux use

```
sudo yum install perl-Switch perl-DateTime perl-Sys-Syslog perl-LWP-Protocol-https -y
```

- At a command prompt, move to a folder where you want to store the monitoring scripts and run the following command to download them:

```
curl https://aws-cloudwatch.s3.amazonaws.com/downloads/CloudWatchMonitoringScripts-1.2.2.zip -O
sudo yum -y install perl-Digest-SHA.x86_64
```

- Run below commands to install the monitoring scripts which we just downloaded

```
unzip CloudWatchMonitoringScripts-1.2.2.zip && \
rm CloudWatchMonitoringScripts-1.2.2.zip && \
cd aws-scripts-mon
```

- Make sure your Instance has permissions to publish metrics in cloudwatch (This can be achieved using an IAM role)
- The following example collects all available memory metrics and sends them to CloudWatch, counting cache and buffer memory as used

```
./mon-put-instance-data.pl --mem-used-incl-cache-buff --mem-util --mem-used --mem-avail
```

- Above script will send data to cloudwatch everytime it is executed
- We can schedule it to send data periodically by using cronjob as well
- in order to set the cronjob use

```
crontab -e
```

- Enter the cron expression as below

```
*/1 * * * * /home/ec2-user/aws-scripts-mon/mon-put-instance-data.pl --mem-used-incl-cache-buff --mem-util --mem-used --mem-avail
```

- save the file and exit
- Above setting will make sure the script is run every minute
- You should be able to see the graph in cloudwatch metrics under custom namespaces . Or you can search using the instance id

## Auto scaling

- Auto scaling is a part of horizontal scaling strategy wherein we scale the instances based on scaling policies
- These scaling policies are derived over any of the cloudwatch metric
- For ex. scale the instances if the cpu utilization goes beyond 50%
- These scaling policies can be configured for scaling in or scaling out actions as well

## Configuring auto scaling group

- Navigate to EC2
- Go to launch configurations - This decides the ec2 instances which will be launched ,should have which configuration
- Select the AMI - Mostly in industry , we use custom AMI where our application is already installed
- Select the instance type

- Under additional configuration we have an option of using spot request . which we can skip for now
- Under IAM role, if we require any access to other service we can create and select the role
- We can enable detailed monitoring as of now as we want to test the scaling results asap
- Select appropriate amount of volume . Can be kept as default
- Configure the security group and keyname as we do it for EC2 and click on create launch configuration
- Once launch configuration is created , we can move to auto scaling group
- Select the launch configuration which was just created and click on create auto scaling group
- Give to auto scaling group a name and click on next

### Choose launch template or configuration

Step 2  
Configure settings

Step 3 (optional)  
Configure advanced options

Step 4 (optional)  
Configure group size and scaling policies

Step 5 (optional)  
Add notifications

Step 6 (optional)  
Add tags

Step 7  
Review

### Choose launch template or configuration Info

Specify a launch template that contains settings common to all EC2 instances that are launched by this Auto Scaling group. If you currently use launch configurations, you might consider migrating to launch templates.

#### Name

Auto Scaling group name  
Enter a name to identify the group.

Must be unique to this account in the current Region and no more than 255 characters.

#### Launch configuration Info

[Switch to launch template](#)

Launch configuration  
Choose a launch configuration that contains the instance-level settings, such as the Amazon Machine Image (AMI), instance type, key pair, and security groups.

[Create a launch configuration ↗](#)

Launch configuration	AMI ID	Date created
LC-1	ami-0947d2ba12ee1ff75	Sat Oct 17 2020 13:32:10 GMT+0530 (India Standard Time)
Security groups	Instance type	Key pair name

- Select the VPC and subnet in which we want our instances to launch . We can select multiple subnets or can stick with one. Click next
- Under advanced options we have the ability to put the auto scaled instances behind an ELB . We can select a target group under which these instances will be registered
- Health checks is based on which the auto scaling group decides if more instances are needed . WE can keep them as default which is EC2 status check . Additionally we can also add elb health check
- Click next
- Under group size we can select the desired, minimum and maximum limit we want our instance count to be
- Let us keep desired as 3 . Minimum as 1 and maximum as 5 . ASG will try and keep the count till 3 . where minimum count will be 1 and maximum will be 5
- For scaling policies ,we can define if we want to dynamically scale our ec2 instances based on metrics. Which we do . Hence select the "target tracking scaling policy "
- Let us keep the scaling policy name as it is . Chose the metric name as Average CPU utilization
- Target value can kept as 50 .
- keep other options as default and click next
- Under notifications we can add notification for each action performed to the SNS topic. This is optional
- Add tags if required and go ahead and create the auto scaling group
- Since we kept the minimum value as 1 , we can see the instance getting launched in the instances tab
- If we need to test the auto scaling however , we need to load the instance in such a way that cpu utilization goes beyond 50%
- In order to acheive that , we have a linux utility called stress . Stress hogs the resources for a certain amount of time and then releases them.
- Log in the instance which has been launched and run below commands

```

sudo amazon-linux-extras install epel -y
sudo yum install stress -y
stress --cpu 1 --timeout 300
stress --vm 4 --vm-bytes 102M

```

- Above commands will install stress utility and hog the resources
- You can monitor the same from the metrics
- Observe that once the CPU utilization graph goes beyond 50 , other instances will be launched as auto scaling group will be triggered