

## Table of contents

- [Table of contents](#)
- [Linux basic operations](#)
  - [User creation](#)
  - [Enabling password based authentication for EC2](#)
  - [Granting sudo privileges to user](#)
  - [Scaling the storage capacity of the instance](#)
  - [Cronjob](#)

## Linux basic operations

- Agenda for today would be understanding and performing some of the basic linux operations that will help our understanding

### User creation

- Every operating system has a default user that comes with it . For ex . Amazon linux 1 and 2 has "ec2-user" as the default user . Ubuntu operating system has "ubuntu" and so on
- We can create other users as well based on the requirement .
- These users are generally created for project teams which use the instance for various development purposes

```
sudo adduser test1
```

- Above command will create a user named test1
- Now that we have created a user , in order to use it , that user need to have a password

```
sudo passwd test1
```

- Once you use above command , it will prompt for entering a password .
- To test the above , login to the machine as ec2-user and use

```
su test1
```

- su stands for switch user . Above command will ask for a password , once it is entered , you will be able to log in as test1 user
- Super users do not require a password while using su command. So from root you can switch to any user that you wish without having to worry about entering the password

## Enabling password based authentication for EC2

- Since the pem key allows access to ec2-user which has sudo access , ideally we shouldnt be sharing the key with the project users
- We should be creating seperate users for each user
- By Default , ec2 instances only have key based authentication . We need to enable password based authentication in order for users to access the instance with just username and password

```
cd /etc/ssh/  
sudo vi sshd_config
```

- Once the file is opened , change the PasswordAuthentication as yes

```
# Authentication:  
  
#LoginGraceTime 2m  
#PermitRootLogin yes  
#StrictModes yes  
#MaxAuthTries 6  
#MaxSessions 10  
  
#PubkeyAuthentication yes  
  
# The default is to check both .ssh/authorized_keys and .ssh/authorized_keys2  
# but this is overridden so installations will only check .ssh/authorized_keys  
AuthorizedKeysFile .ssh/authorized_keys  
  
#AuthorizedPrincipalsFile none  
  
# For this to work you will also need host keys in /etc/ssh/ssh_known_hosts  
#HostbasedAuthentication no  
# Change to yes if you don't trust ~/.ssh/known_hosts for  
# HostbasedAuthentication  
#IgnoreUserKnownHosts no  
# Don't read the user's ~/.rhosts and ~/.shosts files  
#IgnoreRhosts yes  
  
# To disable tunneled clear text passwords, change to no here!  
#PasswordAuthentication yes  
#PermitEmptyPasswords no  
PasswordAuthentication yes  
  
# Change to no to disable s/key passwords  
#ChallengeResponseAuthentication yes  
ChallengeResponseAuthentication no  
  
# Kerberos options  
#KerberosAuthentication no  
-- INSERT --
```

- Once the changes are made save and exit

```
sudo service sshd restart
```

- Above command will restart the ssh daemon which reloads the configuration changes if any
- You can now try logging in using just the username (for ex- test1) and password

## Granting sudo privileges to user

- Whenever any new user is created , it only has permissions to operate on its own home folder i.e. /home/test1
- If there is a scenario where the user needs permissions to install or make some changes in other folder , user will need sudo privileges
- In order to grant the permissions

```
sudo visudo
```

- This will open a file , make the changes as below

```
Defaults    env_keep += "MAIL PS1 PS2 QIDIR USERNAME LANG LC_ADDRESS LC_CTYPE"
Defaults    env_keep += "LC_COLLATE LC_IDENTIFICATION LC_MEASUREMENT LC_MESSAGES"
Defaults    env_keep += "LC_MONETARY LC_NAME LC_NUMERIC LC_PAPER LC_TELEPHONE"
Defaults    env_keep += "LC_TIME LC_ALL LANGUAGE LINGUAS _XKB_CHARSET XAUTHORITY"

#
# Adding HOME to env_keep may enable a user to run unrestricted
# commands via sudo.
#
# Defaults    env_keep += "HOME"

Defaults    secure_path = /sbin:/bin:/usr/sbin:/usr/bin

## Next comes the main part: which users can run what software on
## which machines (the sudoers file can be shared between multiple
## systems).
## Syntax:
##
##     user    MACHINE=COMMANDS
##
## The COMMANDS section may have other options added to it.
##
## Allow root to run any commands anywhere
root        ALL=(ALL)        ALL
test1      ALL=(ALL)        ALL
## Allows members of the 'sys' group to run networking, software,
## service management apps and more.
# %sys ALL = NETWORKING, SOFTWARE, SERVICES, STORAGE, DELEGATING, PROCESSES, LOCATE, DRIVERS

## Allows people in group wheel to run all commands
%wheel      ALL=(ALL)        ALL

## Same thing without a password
# %wheel    ALL=(ALL)        NOPASSWD: ALL

## Allows members of the users group to mount and unmount the
## cdrom as root
# %users    ALL=/sbin/mount /mnt/cdrom, /sbin/umount /mnt/cdrom
```

- Save the file and exit . Post which user will be able to run sudo commands

## Scaling the storage capacity of the instance

- We know that we can change the instance type if we run out of computig capacity . In order to scale the storage capacity there are couple of approaches we can take
  - Scale the existing ebs volume
    - The general purpose ebs volume has upto 16tb of capacity . If a disc runs out of capacity , we can scale the existing volume

```
[ec2-user@ip-172-31-16-9 ~]$ lsblk
NAME        MAJ:MIN RM  SIZE RO TYPE MOUNTPOINT
xvda         202:0    0   8G  0 disk 
└─xvda1      202:1    0   8G  0 part /
[ec2-user@ip-172-31-16-9 ~]$ df -h
Filesystem      Size  Used Avail Use% Mounted on
devtmpfs        482M   0   482M   0% /dev
tmpfs           492M   0   492M   0% /dev/shm
tmpfs           492M 456K   492M   1% /run
tmpfs           492M   0   492M   0% /sys/fs/cgroup
/dev/xvda1      8.0G  1.4G   6.7G  17% /
tmpfs           99M   0    99M   0% /run/user/1000
tmpfs           99M   0    99M   0% /run/user/0
[ec2-user@ip-172-31-16-9 ~]$
```

- As seen above , the root volume is of 8 gb . It is currently 17% used . We will now scale it
- Navigate to the aws console
- Go to the instance , and select the ebs volume attached to it . It navigates to the volumes screen . Click on actions and select modify volume
- Remember that we can only increase ebs volume .
- Add the new capacity and click on modify

**Modify Volume** ✕

**Volume ID** vol-029b2807e3d545239

**Volume Type** General Purpose SSD (gp2) i

**Size**  (Min: 1 GiB, Max: 16384 GiB) i

**IOPS** 100 / 3000 (Baseline of 3 IOPS per GiB with a minimum of 100 IOPS, burstable to 3000 IOPS) i

Cancel Modify

- Once the volume is modified , come back to the terminal

```
[ec2-user@ip-172-31-16-9 ~]$ lsblk
NAME        MAJ:MIN RM SIZE RO TYPE MOUNTPOINT
xvda        202:0    0  8G  0 disk
└─xvda1    202:1    0  8G  0 part /
[ec2-user@ip-172-31-16-9 ~]$ df -h
Filesystem      Size  Used Avail Use% Mounted on
devtmpfs        482M   0  482M   0% /dev
tmpfs           492M   0  492M   0% /dev/shm
tmpfs           492M  456K  492M   1% /run
tmpfs           492M   0  492M   0% /sys/fs/cgroup
/dev/xvda1      8.0G  1.4G  6.7G  17% /
tmpfs           99M    0   99M   0% /run/user/1000
tmpfs           99M    0   99M   0% /run/user/0
[ec2-user@ip-172-31-16-9 ~]$ df -h
Filesystem      Size  Used Avail Use% Mounted on
devtmpfs        482M   0  482M   0% /dev
tmpfs           492M   0  492M   0% /dev/shm
tmpfs           492M  456K  492M   1% /run
tmpfs           492M   0  492M   0% /sys/fs/cgroup
/dev/xvda1      8.0G  1.4G  6.7G  17% /
tmpfs           99M    0   99M   0% /run/user/1000
tmpfs           99M    0   99M   0% /run/user/0
[ec2-user@ip-172-31-16-9 ~]$
[ec2-user@ip-172-31-16-9 ~]$ lsblk
NAME        MAJ:MIN RM SIZE RO TYPE MOUNTPOINT
xvda        202:0    0 20G  0 disk
└─xvda1    202:1    0  8G  0 part /
[ec2-user@ip-172-31-16-9 ~]$
```

- Notice that the volume is increased , but it wont be usable until we extend the filesystem

```
# based on the block id use
sudo growpart /dev/xvda 1
#above command will increase the partition to accomodate the new size
lsblk
#observe the changes from previous output
df -hT
# if you dont see any changes reflected in above command , use
sudo xfs_growfs -d /
df -hT
```

```
[ec2-user@ip-172-31-16-9 ~]$ sudo growpart /dev/xvda 1
CHANGED: partition=1 start=4096 old: size=16773087 end=16777183 new: size=41938911 end=41943007
```

```
[ec2-user@ip-172-31-16-9 ~]$ lsblk
NAME        MAJ:MIN RM SIZE RO TYPE MOUNTPOINT
xvda        202:0    0 20G  0 disk
└─xvda1    202:1    0 20G  0 part /
[ec2-user@ip-172-31-16-9 ~]$ df -h
Filesystem      Size  Used Avail Use% Mounted on
devtmpfs        482M   0  482M   0% /dev
tmpfs           492M   0  492M   0% /dev/shm
tmpfs           492M  464K  492M   1% /run
tmpfs           492M   0  492M   0% /sys/fs/cgroup
/dev/xvda1      8.0G  1.4G  6.7G  17% /
tmpfs           99M    0   99M   0% /run/user/1000
tmpfs           99M    0   99M   0% /run/user/0
[ec2-user@ip-172-31-16-9 ~]$ sudo xfs_growfs -d /
meta-data=/dev/xvda1            isize=512    agcount=4, agsize=524159 blks
=                               sectsz=512   attr=2, projid32bit=1
=                               crc=1      finobt=1 spinodes=0
data      =                       bsize=4096 blocks=2096635, imaxpct=25
=                               sunit=0     swidth=0 blks
naming    =version 2           bsize=4096 ascii-ci=0 ftype=1
log       =internal          bsize=4096 blocks=2560, version=2
=                               sectsz=512   sunit=0 blks, lazy-count=1
realtime  =none              extsz=4096  blocks=0, rtextents=0
data blocks changed from 2096635 to 5242363
[ec2-user@ip-172-31-16-9 ~]$ df -hT
Filesystem      Type      Size  Used Avail Use% Mounted on
devtmpfs        devtmpfs  482M   0  482M   0% /dev
tmpfs           tmpfs     492M   0  492M   0% /dev/shm
tmpfs           tmpfs     492M  464K  492M   1% /run
tmpfs           tmpfs     492M   0  492M   0% /sys/fs/cgroup
/dev/xvda1      xfs       20G    1.4G  19G    7% /
tmpfs           tmpfs     99M    0   99M   0% /run/user/1000
tmpfs           tmpfs     99M    0   99M   0% /run/user/0
[ec2-user@ip-172-31-16-9 ~]$
```

- This way we can scale an existing volume
- Since the general purpose ebs has 16tb limit , some teams will run out of storage .
- In that scenario we can create new ebs volume and attach it to the instance
- In order to do that , first we need to create a volume
- Navigate to the volumes screen in ec2 service
- Click on create new volume
- Make sure this volume is created in the exact same AZ as the instance
- Once this volume is created attach it the instance
- Notice the changes by using "lsblk" command

```
[ec2-user@ip-172-31-16-9 ~]$ lsblk
NAME        MAJ:MIN RM  SIZE RO TYPE MOUNTPOINT
xvda        202:0    0   20G  0 disk
└─xvda1     202:1    0   20G  0 part /
[ec2-user@ip-172-31-16-9 ~]$ lsblk
NAME        MAJ:MIN RM  SIZE RO TYPE MOUNTPOINT
xvda        202:0    0   20G  0 disk
└─xvda1     202:1    0   20G  0 part /
xvdf        202:80    0  100G  0 disk
[ec2-user@ip-172-31-16-9 ~]$
```

- Notice that the mountpoint section in front of newly added block device is blank . Which means it is not usable yet
- First we will have to create a file system on the volume . Remember , without file system we cannot store data on any disc

```
# first let us check if the volume has a file system
sudo file -s /dev/xvdf
# If output of above command comes as "data", then it means it does not have a
file system and we can go ahead and create it .
# However if the output is x86 or xfs , it means it already has a filesystem and
we shoudnt create it . If you create a filesystem on this it will erase all the
existing data . SO make sure to check before you proceed
sudo mkfs -t xfs /dev/xvdf
# Above command will create a file system on the disc
#now we need a directory where the disc needs to be mounted. It is basically the
address where the disc will be utilized
sudo mkdir /data
sudo mount /dev/xvdf /data
```

```
[ec2-user@ip-172-31-16-9 ~]$ sudo file -s /dev/xvdf
/dev/xvdf: data
```

```
[ec2-user@ip-172-31-16-9 ~]$ sudo mkfs -t x86 /dev/xvdf
mkfs: failed to execute mkfs.x86: No such file or directory
[ec2-user@ip-172-31-16-9 ~]$ sudo mkfs -t xfs /dev/xvdf
meta-data=/dev/xvdf          isize=512    agcount=4, agsize=6553600 blks
=                           sectsz=512   attr=2, projid32bit=1
=                           crc=1        finobt=1, sparse=0
data      =                   bsize=4096   blocks=26214400, imaxpct=25
=                           sunit=0      swidth=0 blks
naming    =version 2          bsize=4096   ascii-ci=0 ftype=1
log       =internal log      bsize=4096   blocks=12800, version=2
=                           sectsz=512   sunit=0 blks, lazy-count=1
realtime  =none              extsz=4096   blocks=0, rtextents=0
[ec2-user@ip-172-31-16-9 ~]$
```

```
[ec2-user@ip-172-31-16-9 ~]$ sudo mkdir /data
[ec2-user@ip-172-31-16-9 ~]$ sudo mount /dev/xvdf /data
[ec2-user@ip-172-31-16-9 ~]$ lsblk
NAME        MAJ:MIN RM  SIZE RO TYPE MOUNTPOINT
xvda        202:0    0   20G  0 disk
└─xvda1     202:1    0   20G  0 part /
xvdf        202:80   0  100G  0 disk /data
[ec2-user@ip-172-31-16-9 ~]$
```

- Now the directory and the new volume are usable
- This mounted directory however will be reset if the instance goes through a reboot
- In order to avoid that , we need to enter these new changes in fstab file
- The entry requires a UUID , which is globally unique

`sudo blkid`

# above command will give you UUID of all the block devices attached to the machine. Copy for the newly created volume

# Now we will make the changes the fstab file. fstab file is one of the most crucial files in linux, if there are any mistakes in the entry , then after the next reboot , you will not be able to login on the instance . So make the changes carefully , you can also take backup of the current config by creating a copy of it

`sudo vi /etc/fstab`

#save and exit

# in order to confirm if all the entries in fstab are correct try below command

`sudo mount -a`

# If above command runs successfully , then the entry syntax is correct . You can try rebooting the instance and verify the same

```
[ec2-user@ip-172-31-16-9 ~]$ sudo cat /etc/fstab
#
UUID=15c7809d-e6e3-4062-a5eb-afeb1939fc6e    /          xfs     defaults,noatime 1 1
[ec2-user@ip-172-31-16-9 ~]$ sudo blkid
/dev/xvda1: LABEL="/" UUID="15c7809d-e6e3-4062-a5eb-afeb1939fc6e" TYPE="xfs" PARTLABEL="Linux" PARTUUID="cfb2e895-ecec-41c5-afc2-40ffe2e4384a"
/dev/xvdf: UUID="434aa9a1-f68c-4254-af5a-lbc8df9f35f3" TYPE="xfs"
[ec2-user@ip-172-31-16-9 ~]$ sudo vi /etc/fstab
[ec2-user@ip-172-31-16-9 ~]$
```

```
#
UUID=15c7809d-e6e3-4062-a5eb-afeb1939fc6e    /          xfs     defaults,noatime 1 1
UUID=434aa9a1-f68c-4254-af5a-lbc8df9f35f3    /data      xfs     defaults,noatime 1 1
~
~
~
~
~
~
~
~
```

## Cronjob

- Cronjobs are a great way of scheduling linux level tasks or scripts
- We have briefly seen cloudwatch rules which used cron expression
- For ex we want to run "uptime " command which shows duration since the last reboot
- uptime >> test.txt will store the output in the test.txt file

```
[ec2-user@ip-172-31-16-9 ~]$ uptime
20:22:21 up 1:37, 1 user, load average: 0.00, 0.00, 0.00
[ec2-user@ip-172-31-16-9 ~]$ uptime >> test.txt
[ec2-user@ip-172-31-16-9 ~]$ cat test.txt
20:22:32 up 1:37, 1 user, load average: 0.00, 0.00, 0.00
[ec2-user@ip-172-31-16-9 ~]$
```

- If we wish to schedule above command to run at a certain schedule , we can use cronjob

```
crontab -e
```

```
# this opens up the cronjob editor which is exactly like vi editor
# For testing purposes we will make entry "* * * * * uptime >> test.txt"
# Save and exit
```

```
crontab -l
```

```
# use above to list the cronjobs. Remember cronjobs are user specific , so while scheduling dont use sudo or else the jobs will be scheduled from root user
```

```
* * * * * uptime >> test.txt
```

```
[ec2-user@ip-172-31-16-9 ~]$ cat test.txt
20:22:32 up 1:37, 1 user, load average: 0.00, 0.00, 0.00
20:27:01 up 1:42, 1 user, load average: 0.00, 0.00, 0.00
20:28:01 up 1:43, 1 user, load average: 0.00, 0.00, 0.00
[ec2-user@ip-172-31-16-9 ~]$
```



- This can even be used to schedule shell/python scripts to be executed at a certain schedule