

Table of Contents

- [Table of Contents](#)
 - [What is EKS](#)
 - [Steps to create EKS](#)
 - [EKSCTL CLI](#)
 - [Verify Cluster, NodeGroup in EKS Management Console](#)
 - [List Worker Nodes](#)
 - [kubectl commands](#)
 - [Reference](#)
 - [EKS Cluster Pricing](#)
 - [EKS Cluster Pricing](#)
 - [EKS Worker Nodes Pricing-EC2](#)

What is EKS

- AWS EKS is a **Managed Kubernetes Service** from Amazon, which means AWS manages the Master Nodes.
- All the necessary applications/services are already pre-installed like the container runtime or master processes and in addition it also takes care of scaling and backups.
- You only create the Worker Nodes.

Steps to create EKS

- To create a K8s cluster in EKS you need to do following steps:
- Create a VPC
- Create an IAM role with Security Group (or in other words: create AWS user with list of permissions)
- Create Cluster Control Plane - Master Nodes
 - choose basic information like cluster name and k8s version
 - choose region and VPC for your cluster
 - set security
- Create Worker Nodes and connect to cluster
 - The Worker Nodes are some EC2 instances with CPU and storage resources.
 - Create as a Node Group
 - Choose cluster it will attach to
 - Define Security Group, select instance type etc.

With **NodeGroup** you have autoscaling, depending on how much load the cluster has new Worker Nodes will automatically added or removed in the cluster.

EKSCTL CLI

- Execute below commands in either one of the below options:
 - AWS CloudShell (Uses AWS Console Login IAM Credentials)
 - Local Linux Machine (Uses IAM User Credentials configured)
 - Amazon EC2 Linux Machine (Uses AWS IAM Role Permissions)

- Download the `eksctl` on your machine linux cli

```
curl --silent --location
"https://github.com/weaveworks/eksctl/releases/latest/download/eksctl_$(uname -
s)_amd64.tar.gz" | tar xz -C /tmp
sudo mv /tmp/eksctl /usr/local/bin
```

- The IAM account used for EKS cluster creation should have these minimal access levels.

AWS Service	Access Level
CloudFormation	Full Access
EC2	Full: Tagging Limited: List, Read, Write
EC2 Auto Scaling	Limited: List, Write
EKS	Full Access
IAM	Limited: List, Read, Write, Permissions Management
Systems Manager	Limited: List, Read

- Command to create EKS Cluster

```
eksctl create cluster \
--name test-eks-cluster \
--version 1.21 \
--region ap-south-1 \
--nodegroup-name eks-worker-nodegroup \
--node-type t2.micro \
--nodes 2

#####OUTPUT#####
2022-01-26 01:53:53 [ i ] eksctl version 0.80.0
2022-01-26 01:53:53 [ i ] using region ap-south-1
2022-01-26 01:53:54 [ i ] setting availability zones to [ap-south-1b ap-south-1a
ap-south-1c]
2022-01-26 01:53:54 [ i ] subnets for ap-south-1b - public:192.168.0.0/19
private:192.168.96.0/19
2022-01-26 01:53:54 [ i ] subnets for ap-south-1a - public:192.168.32.0/19
private:192.168.128.0/19
2022-01-26 01:53:54 [ i ] subnets for ap-south-1c - public:192.168.64.0/19
private:192.168.160.0/19
2022-01-26 01:53:54 [ i ] nodegroup "ng-2564739d" will use "" [AmazonLinux2/1.21]
2022-01-26 01:53:54 [ i ] using Kubernetes version 1.21
2022-01-26 01:53:54 [ i ] creating EKS cluster "test-eks-cluster" in "ap-south-1"
region with managed nodes
2022-01-26 01:53:54 [ i ] will create 2 separate CloudFormation stacks for cluster
itself and the initial managed nodegroup
2022-01-26 01:53:54 [ i ] if you encounter any issues, check CloudFormation
console or try 'eksctl utils describe-stacks --region=ap-south-1 --cluster=test-
```

```
eks-cluster'
2022-01-26 01:53:54 [ i ]   Kubernetes API endpoint access will use default of
{publicAccess=true, privateAccess=false} for cluster "test-eks-cluster" in "ap-
south-1"
2022-01-26 01:53:54 [ i ]   CloudWatch logging will not be enabled for cluster
"test-eks-cluster" in "ap-south-1"
2022-01-26 01:53:54 [ i ]   you can enable it with 'eksctl utils update-cluster-
logging --enable-types={SPECIFY-YOUR-LOG-TYPES-HERE (e.g. all)} --region=ap-south-
1 --cluster=test-eks-cluster'
2022-01-26 01:53:54 [ i ]
2 sequential tasks: { create cluster control plane "test-eks-cluster",
  2 sequential sub-tasks: {
    wait for control plane to become ready,
    create managed nodegroup "ng-2564739d",
  }
}
2022-01-26 01:53:54 [ i ]   building cluster stack "eksctl-test-eks-cluster-cluster"
2022-01-26 01:53:55 [ i ]   deploying stack "eksctl-test-eks-cluster-cluster"
2022-01-26 01:54:25 [ i ]   waiting for CloudFormation stack "eksctl-test-eks-
cluster-cluster"
#####OUTPUT#####
```

- Install `kubect1` linux utility on your local linux client.

```
curl -o kubect1 https://amazon-eks.s3.us-west-2.amazonaws.com/1.21.2/2021-07-
05/bin/linux/amd64/kubect1
chmod +x ./kubect1
mkdir -p $HOME/bin && cp ./kubect1 $HOME/bin/kubect1 && export
PATH=$PATH:$HOME/bin
echo 'export PATH=$PATH:$HOME/bin' >> ~/.bashrc
```

```
# View the Local kubeconfig file
kubect1 config view
# This command shows content inside the ~/.kube/config file
#####OUTPUT#####
apiVersion: v1
clusters:
- cluster:
  certificate-authority-data: DATA+OMITTED
  server: https://00D6914D1CB90D39116178B3A9A2ECA6.y14.ap-south-
1.eks.amazonaws.com
  name: test-eks-cluster.ap-south-1.eksctl.io
contexts:
- context:
  cluster: test-eks-cluster.ap-south-1.eksctl.io
  user: ansible-user@test-eks-cluster.ap-south-1.eksctl.io
  name: ansible-user@test-eks-cluster.ap-south-1.eksctl.io
current-context: ansible-user@test-eks-cluster.ap-south-1.eksctl.io
```

```

kind: Config
preferences: {}
users:
- name: ansible-user@test-eks-cluster.ap-south-1.eksctl.io
  user:
    exec:
      apiVersion: client.authentication.k8s.io/v1alpha1
      args:
        - eks
        - get-token
        - --cluster-name
        - test-eks-cluster
        - --region
        - ap-south-1
      command: aws
      env:
        - name: AWS_STS_REGIONAL_ENDPOINTS
          value: regional
      provideClusterInfo: false
#####OUTPUT#####

```

When you create an Amazon EKS cluster, the AWS Identity and Access Management (IAM) entity user or role that creates the cluster, is automatically granted **system:masters** permissions in the cluster's **role-based access control** (RBAC) configuration in the Amazon EKS control plane. This IAM entity doesn't appear in any visible configuration, so make sure to keep track of which IAM entity originally created the cluster.

- Validate the CloudFormation Template Resources created with above command.

Verify Cluster, NodeGroup in EKS Management Console

- Go to Services -> Elastic Kubernetes Service

List Worker Nodes

```

# List EKS clusters
eksctl get cluster

# List NodeGroups in a cluster
eksctl get nodegroup --cluster=<clusterName>

# List Nodes in current kubernetes cluster
kubectl get nodes -o wide

# Our kubectl context should be automatically changed to new cluster
kubectl config view --minify

```

kubectl commands

```
kubectl get pods -n kube-system
```

- Create the Deployment by running the following command:

```
kubectl apply -f https://k8s.io/examples/controllers/nginx-deployment.yaml
```

```
[cloudshell-user@ip-10-0-50-112 ~]$ kubectl get events
```

LAST SEEN	TYPE	REASON	OBJECT
MESSAGE			
2m48s	Normal	Scheduled	pod/nginx-deployment-66b6c48dd5-qkst9
Successfully assigned default/nginx-deployment-66b6c48dd5-qkst9 to ip-192-168-21-139.ap-south-1.compute.internal			
2m47s	Normal	Pulling	pod/nginx-deployment-66b6c48dd5-qkst9
Pulling image "nginx:1.14.2"			
2m38s	Normal	Pulled	pod/nginx-deployment-66b6c48dd5-qkst9
Successfully pulled image "nginx:1.14.2" in 9.12123707s			
2m37s	Normal	Created	pod/nginx-deployment-66b6c48dd5-qkst9
Created container nginx			
2m37s	Normal	Started	pod/nginx-deployment-66b6c48dd5-qkst9
Started container nginx			
37s	Warning	FailedScheduling	pod/nginx-deployment-66b6c48dd5-x92pj
0/2 nodes are available: 2 Too many pods.			
2m48s	Normal	Scheduled	pod/nginx-deployment-66b6c48dd5-zb9bb
Successfully assigned default/nginx-deployment-66b6c48dd5-zb9bb to ip-192-168-46-56.ap-south-1.compute.internal			
2m47s	Normal	Pulling	pod/nginx-deployment-66b6c48dd5-zb9bb
Pulling image "nginx:1.14.2"			
2m37s	Normal	Pulled	pod/nginx-deployment-66b6c48dd5-zb9bb
Successfully pulled image "nginx:1.14.2" in 9.470360458s			
2m37s	Normal	Created	pod/nginx-deployment-66b6c48dd5-zb9bb
Created container nginx			
2m37s	Normal	Started	pod/nginx-deployment-66b6c48dd5-zb9bb
Started container nginx			
2m48s	Normal	SuccessfulCreate	replicaset/nginx-deployment-66b6c48dd5
Created pod: nginx-deployment-66b6c48dd5-zb9bb			
2m48s	Normal	SuccessfulCreate	replicaset/nginx-deployment-66b6c48dd5
Created pod: nginx-deployment-66b6c48dd5-qkst9			
2m48s	Normal	SuccessfulCreate	replicaset/nginx-deployment-66b6c48dd5
Created pod: nginx-deployment-66b6c48dd5-x92pj			
2m48s	Normal	ScalingReplicaSet	deployment/nginx-deployment
Scaled up replica set nginx-deployment-66b6c48dd5 to 3			

```
kubectl describe pods -n default > pods.txt
```

```
cat pods.txt
```

```
Name:          nginx-deployment-66b6c48dd5-qkst9
Namespace:     default
Priority:       0
Node:          ip-192-168-21-139.ap-south-1.compute.internal/192.168.21.139
Start Time:    Wed, 02 Feb 2022 19:29:23 +0000
Labels:        app=nginx
```

```

    pod-template-hash=66b6c48dd5
Annotations:  kubernetes.io/psp: eks.privileged
Status:      Running
IP:          192.168.15.247
IPs:
  IP:        192.168.15.247
Controlled By:  ReplicaSet/nginx-deployment-66b6c48dd5
Containers:
  nginx:
    Container ID:
docker://cb8aaf025035aa3a81d05f1598efbfc945bbe4ae01558a0ee0dee02942f32175
    Image:      nginx:1.14.2
    Image ID:   docker-
pullable://nginx@sha256:f7988fb6c02e0ce69257d9bd9cf37ae20a60f1df7563c3a2a6abe24160
306b8d
    Port:       80/TCP
    Host Port:  0/TCP
    State:      Running
      Started:  Wed, 02 Feb 2022 19:29:34 +0000
    Ready:      True
    Restart Count:  0
    Environment:  <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-5qrpp
(ro)
Conditions:
  Type            Status
  Initialized      True
  Ready           True
  ContainersReady True
  PodScheduled    True
Volumes:
  kube-api-access-5qrpp:
    Type:      Projected (a volume that contains injected data from
multiple sources)
    TokenExpirationSeconds: 3607
    ConfigMapName:         kube-root-ca.crt
    ConfigMapOptional:    <nil>
    DownwardAPI:          true
QoS Class:           BestEffort
Node-Selectors:      <none>
Tolerations:         node.kubernetes.io/not-ready:NoExecute op=Exists for
300s
                    node.kubernetes.io/unreachable:NoExecute op=Exists
for 300s
Events:
  Type    Reason      Age   From          Message
  ----    -
  Normal  Scheduled   9m3s  default-scheduler  Successfully assigned
default/nginx-deployment-66b6c48dd5-qkst9 to ip-192-168-21-139.ap-south-
1.compute.internal
  Normal  Pulling     9m2s  kubelet        Pulling image "nginx:1.14.2"
  Normal  Pulled      8m53s  kubelet        Successfully pulled image
"nginx:1.14.2" in 9.12123707s

```

Normal	Created	8m52s	kubelet	Created container nginx
Normal	Started	8m52s	kubelet	Started container nginx

Name: nginx-deployment-66b6c48dd5-x92pj
Namespace: default
Priority: 0
Node: <none>
Labels: app=nginx
pod-template-hash=66b6c48dd5
Annotations: kubernetes.io/psp: eks.privileged
Status: Pending
IP: <none>
IPs: <none>
Controlled By: ReplicaSet/nginx-deployment-66b6c48dd5
Containers:
 nginx:
 Image: nginx:1.14.2
 Port: 80/TCP
 Host Port: 0/TCP
 Environment: <none>
 Mounts:
 /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-x7mxn (ro)
Conditions:
 Type Status
 PodScheduled False
Volumes:
 kube-api-access-x7mxn:
 Type: Projected (a volume that contains injected data from multiple sources)
 TokenExpirationSeconds: 3607
 ConfigMapName: kube-root-ca.crt
 ConfigMapOptional: <nil>
 DownwardAPI: true
QoS Class: BestEffort
Node-Selectors: <none>
Tolerations: node.kubernetes.io/not-ready:NoExecute op=Exists for 300s
 node.kubernetes.io/unreachable:NoExecute op=Exists for 300s
Events:
 Type Reason Age From Message
 ---- -----
 Warning FailedScheduling 52s (x9 over 9m3s) default-scheduler 0/2 nodes are available: 2 Too many pods.

Name: nginx-deployment-66b6c48dd5-zb9bb
Namespace: default
Priority: 0
Node: ip-192-168-46-56.ap-south-1.compute.internal/192.168.46.56
Start Time: Wed, 02 Feb 2022 19:29:23 +0000
Labels: app=nginx

```

    pod-template-hash=66b6c48dd5
Annotations:  kubernetes.io/psp: eks.privileged
Status:      Running
IP:          192.168.51.252
IPs:
  IP:        192.168.51.252
Controlled By:  ReplicaSet/nginx-deployment-66b6c48dd5
Containers:
  nginx:
    Container ID:
docker://9558a069050aba846aa62160509592ebb85dd82d1d90504bc5b03665a0259d47
    Image:      nginx:1.14.2
    Image ID:    docker-
pullable://nginx@sha256:f7988fb6c02e0ce69257d9bd9cf37ae20a60f1df7563c3a2a6abe24160
306b8d
    Port:       80/TCP
    Host Port:  0/TCP
    State:      Running
      Started:  Wed, 02 Feb 2022 19:29:34 +0000
    Ready:      True
    Restart Count:  0
    Environment:  <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-tnsxj
(ro)
Conditions:
  Type            Status
  Initialized      True
  Ready            True
  ContainersReady  True
  PodScheduled     True
Volumes:
  kube-api-access-tnsxj:
    Type:          Projected (a volume that contains injected data from
multiple sources)
    TokenExpirationSeconds:  3607
    ConfigMapName:          kube-root-ca.crt
    ConfigMapOptional:      <nil>
    DownwardAPI:            true
QoS Class:           BestEffort
Node-Selectors:      <none>
Tolerations:         node.kubernetes.io/not-ready:NoExecute op=Exists for
300s
                     node.kubernetes.io/unreachable:NoExecute op=Exists
for 300s
Events:
  Type    Reason      Age   From          Message
  ----    -
  Normal  Scheduled   9m3s  default-scheduler  Successfully assigned
default/nginx-deployment-66b6c48dd5-zb9bb to ip-192-168-46-56.ap-south-
1.compute.internal
  Normal  Pulling     9m2s  kubelet        Pulling image "nginx:1.14.2"
  Normal  Pulled      8m52s  kubelet        Successfully pulled image
"nginx:1.14.2" in 9.470360458s

```


Normal	Created	8m52s	kubelet	Created container nginx
Normal	Started	8m52s	kubelet	Started container nginx

- The above command shows pods details where the pod gets an IP from ENI in the Subnet.

Network interfaces (6) Info

Filter network interfaces

search: eks Clear filters

Subnet ID	VPC ID	Availability Zone	Interface Type	Description	Instance ID	Status	Public IPv4...	Primary private IPv4...	Secondary private IPv4 add...
subnet-067941fe3a2fce5ab	vpc-02b76a6ce27d533a6	ap-south-1c	Elastic network interface	Amazon EKS test-eks-...	-	In-use	-	192.168.165.131	-
subnet-00f08a720160ae9c2	vpc-02b76a6ce27d533a6	ap-south-1a	Elastic network interface	aws-K8S-i-08266ca7c...	i-08266ca7c22b644a5	In-use	-	192.168.30.4	192.168.15.247
subnet-00f08a720160ae9c2	vpc-02b76a6ce27d533a6	ap-south-1a	Elastic network interface	-	i-08266ca7c22b644a5	In-use	65.0.129.57	192.168.21.139	192.168.7.195
subnet-09aae16454bffc88b	vpc-02b76a6ce27d533a6	ap-south-1b	Elastic network interface	aws-K8S-i-0c544c004...	i-0c544c004ed1ea632	In-use	-	192.168.49.221	192.168.51.252
subnet-004a60aba0384c11a	vpc-02b76a6ce27d533a6	ap-south-1b	Elastic network interface	Amazon EKS test-eks-...	-	In-use	-	192.168.150.81	-
subnet-09aae16454bffc88b	vpc-02b76a6ce27d533a6	ap-south-1b	Elastic network interface	-	i-0c544c004ed1ea632	In-use	3.111.34.130	192.168.46.56	192.168.46.207

- AWS EKS supports native VPC networking with the Amazon VPC Container Network Interface (CNI) plugin for Kubernetes.
- Using this plugin allows Kubernetes Pods to have the same IP address inside the pod as they do on the VPC network.
- For more information, see amazon-vpc-cni-k8s and Proposal: CNI plugin for Kubernetes networking over AWS VPC on GitHub.
- The Amazon VPC CNI plugin is fully supported for use on Amazon EKS and self-managed Kubernetes clusters on AWS.
- Refer the ENI Limit as per Instance Type :
<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/using-eni.html#AvailableIpPerENI>

- The formula for defining the maximum number of Pods per EC2 Node instance is as follows:

$$N * (M-1) + 2$$

- N is the number of Elastic Network Interfaces (ENI) of the instance type

- M is the number of IP addresses per ENI.

- For e.g. **t2.micro** instance, this calculation is $2 * (2-1) + 2 = 4$ Pods

- Use below command to get details on same:

```
aws ec2 describe-instance-types --filters "Name=instance-type,Values=t2.*" --query "InstanceTypes[].[Type: InstanceType, MaxENI: NetworkInfo.MaximumNetworkInterfaces, IPv4addr: NetworkInfo.Ipv4AddressesPerInterface]" --output table
```

DescribeInstanceTypes		
IPv4addr	MaxENI	Type
15	3	t2.2xlarge
2	2	t2.micro
12	3	t2.large
15	3	t2.xlarge
4	3	t2.small
2	2	t2.nano
6	3	t2.medium

```
+-----+-----+-----+-----+
```

```
kubectl get pods --all-namespaces -o wide
kubectl get pods --all-namespaces -o wide | grep -i running
[cloudshell-user@ip-10-0-50-112 ~]$ kubectl get pods --all-namespaces -o wide |
grep -i running
default          nginx-deployment-66b6c48dd5-qkst9  1/1      Running   0   38m
192.168.15.247 ip-192-168-21-139.ap-south-1.compute.internal
default          nginx-deployment-66b6c48dd5-zb9bb  1/1      Running   0   38m
192.168.51.252 ip-192-168-46-56.ap-south-1.compute.internal
kube-system      aws-node-jls58                     1/1      Running   0   125m
192.168.46.56 ip-192-168-46-56.ap-south-1.compute.internal
kube-system      aws-node-tmw12                     1/1      Running   0   125m
192.168.21.139 ip-192-168-21-139.ap-south-1.compute.internal
kube-system      coredns-7f95bc96cc-dl6dn          1/1      Running   0   135m
192.168.46.207 ip-192-168-46-56.ap-south-1.compute.internal
kube-system      coredns-7f95bc96cc-fj56b          1/1      Running   0   135m
192.168.7.193 ip-192-168-21-139.ap-south-1.compute.internal
kube-system      kube-proxy-dqlph                  1/1      Running   0   125m
192.168.21.139 ip-192-168-21-139.ap-south-1.compute.internal
kube-system      kube-proxy-xlpfk                  1/1      Running   0   125m
192.168.46.56 ip-192-168-46-56.ap-south-1.compute.internal
[cloudshell-user@ip-10-0-50-112 ~]$
```

- Delete the cluster using

```
eksctl delete cluster --name test-eks-cluster
```

Reference

- [EKS Cost Optimization](#)
- [Kubernetes instance calculator](#)

EKS Cluster Pricing

- EKS is not free (Unlike other AWS Services), In short, no free-tier for EKS.

EKS Cluster Pricing

- Below is the EKS Cluster Pricing:
 - \$0.10/hour
 - \$2.4/day
 - \$72/month

EKS Worker Nodes Pricing-EC2

- You pay for AWS resources (e.g. EC2 instances or EBS volumes)
- T3 Medium Server in N.Virginia
 - \$0.0416 per Hour

- Per Day: \$0.9984 - Approximately \$1
 - Per Month: \$30 per 1 t3.medium server
- Reference: <https://aws.amazon.com/ec2/pricing/on-demand/>